Gemini 2.5 Flash Image (aka Nano Banana) is now available in the Gemini API!
Learn more

Gemini Developer API Pricing

On this page expand more

Gemini 2.5 Pro
Gemini 2.5 Flash
Gemini 2.5 Flash-Lite
Gemini 2.5 Flash Native Audio
Gemini 2.5 Flash Image Preview
more_horiz

The Gemini API "free tier" is offered through the API service with lower rate limits for testing purposes. Google AI Studio usage is completely free in all available countries.

Standard Batch

	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Free of charge	$1.25, prompts <= 200k tokens $2.50, prompts > 200k tokens
Output price (including thinking tokens)	Free of charge	$10.00, prompts <= 200k tokens $15.00, prompts > 200k
Context caching price	Not available	$0.31, prompts <= 200k tokens $0.625, prompts > 200k $4.50 / 1,000,000 tokens per hour (storage price)
Grounding with Google Search	Not available	1,500 RPD (free), then $35 / 1,000 requests
Used to improve our products	Yes	No

		$1.00 / 1,000,000 tokens per hour (storage price)
Grounding with Google Search	Free of charge, up to 500 RPD (limit shared with Flash-Lite RPD)	1,500 RPD (free, limit shared with Flash-Lite RPD), then $35 / 1,000 requests
Live API	Free of charge	Input: $0.50 (text), $3.00 (audio / image [video]) Output: $2.00 (text), $12.00 (audio)
Used to improve our products	Yes	No

Gemini 2.5 Flash-Lite

Try it in Google AI Studio

Our smallest and most cost effective model, built for at scale usage.

Standard Batch

		Paid Tier, per 1M tokens in

Gemini 2.5 Flash Native Audio

Try it in Google AI Studio

Our native audio models optimized for higher quality audio outputs with better pacing, voice naturalness, verbosity, and mood.

Preview models may change before becoming stable and have more restrictive rate limits.

	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Not available	$0.50 (text) $3.00 (audio / video)
Output price (including thinking tokens)	Not available	$2.00 (text) $12.00 (audio)
Used to improve our products	Yes	No

Gemini 2.5 Flash Image Preview

Try it in Google AI Studio

Gemini 2.5 Flash Preview TTS

Try it in Google AI Studio

Our 2.5 Flash text-to-speech audio model optimized for price-performant, low-latency, controllable speech generation.

Preview models may change before becoming stable and have more restrictive rate limits.

Standard Batch

	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Free of charge	$0.50 (text)
Output price	Free of charge	$10.00 (audio)
Used to improve our products	Yes	No

Used to improve our products	Yes	No

Gemini 2.0 Flash

Try it in Google AI Studio

Our most balanced multimodal model with great performance across all tasks, with a 1 million token context window, and built for the era of Agents.

Standard Batch

	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Free of charge	$0.10 (text / image / video) $0.70 (audio)
Output price	Free of charge	$0.40
Context caching price	Free of charge	$0.025 / 1,000,000 tokens (text/image/video) $0.175 / 1,000,000 tokens (audio)

Gemini 2.0 Flash-Lite

Try it in Google AI Studio

Our smallest and most cost effective model, built for at scale usage.

Standard Batch

	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Free of charge	$0.075
Output price	Free of charge	$0.30
Context caching price	Not available	Not available
Context caching (storage)	Not available	Not available
Tuning price	Not available	Not available
Grounding with Google Search	Not available	Not available

Imagen 4 Standard image price	Not available	$0.04
Imagen 4 Ultra image price	Not available	$0.06
Used to improve our products	Yes	No

Imagen 3

Try it in Google AI Studio

Our state-of-the-art image generation model, available to developers on the paid tier of the Gemini API.

	Free Tier	Paid Tier, per Image in USD
Image price	Not available	$0.03
Used to improve our products	Yes	No

A faster, more accessible version of Veo 3, optimized for speed and business use cases. Available to developers on the paid tier of the Gemini API.

	Free Tier	Paid Tier, per second in USD
Video with audio price (default)	Not available	$0.15
Used to improve our products	Yes	No

star **Note:** In some cases, an audio processing issue may prevent a video from being generated. You will only be charged if your video is successfully generated.

Veo 2

Try the API

Our state-of-the-art video generation model, available to developers on the paid tier of the Gemini API.

Input price	Free of charge	$0.15
Used to improve our products	Yes	No

Gemma 3

Try Gemma 3

Our lightweight, state-of the art, open model built from the same technology that powers our Gemini models.

	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Free of charge	Not available
Output price	Free of charge	Not available
Context caching price	Free of charge	Not available
Context caching (storage)	Free of charge	Not available

Output price	Free of charge	Not available
Context caching price	Free of charge	Not available
Context caching (storage)	Free of charge	Not available
Tuning price	Not available	Not available
Grounding with Google Search	Not available	Not available
Used to improve our products	Yes	No

Gemini 1.5 Flash

Try it in Google AI Studio

Our fastest multimodal model with great performance for diverse, repetitive tasks and a 1 million token context window.

	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Free of charge	$0.075, prompts <= 128k

Gemini 1.5 Flash-8B

Try it in Google AI Studio

Our smallest model for lower intelligence use cases, with a 1 million token context window.

	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Free of charge	$0.0375, prompts <= 128k tokens $0.075, prompts > 128k tokens
Output price	Free of charge	$0.15, prompts <= 128k tokens $0.30, prompts > 128k tokens
Context caching price	Free of charge, up to 1 million tokens of storage per hour	$0.01, prompts <= 128k tokens $0.02, prompts > 128k tokens
Context caching (storage)	Free of charge	$0.25 per hour
Tuning price	Token prices are the same for tuned models	Token prices are the same for tuned models

Input price	Free of charge	$1.25, prompts <= 128k tokens $2.50, prompts > 128k tokens
Output price	Free of charge	$5.00, prompts <= 128k tokens $10.00, prompts > 128k tokens
Context caching price	Not available	$0.3125, prompts <= 128k tokens $0.625, prompts > 128k tokens
Context caching (storage)	Not available	$4.50 per hour
Tuning price	Not available	Not available
Grounding with Google Search	Not available	$35 / 1K grounding requests
Used to improve our products	Yes	No

[*] Google AI Studio usage is free of charge in all available regions. See Billing FAQs for details.

[**] Prices may differ from the prices listed here and the prices offered on Vertex AI. For Vertex prices, see the Vertex AI pricing page.

[***] If you are using dynamic retrieval to optimize costs, only requests that contain at least one grounding support URL from the web in their response are charged for Grounding with Google Search. Costs for Gemini always apply. Rate limits are subject to change.

Except as otherwise noted, the content of this page is licensed under the Creative Commons Attribution 4.0 License, and code samples are licensed under the Apache 2.0 License. For details, see the Google Developers Site Policies. Java is a registered trademark of Oracle and/or its affiliates.

Last updated 2025-09-10 UTC.